SHWETA PAWAR

Junior AI/ML Engineer — LLM, Agentic AI & Computer Vision

Navi Mumbai, Maharashtra | 9076450205 | pawarshweta135@gmail.com | linkedin.com/in/shwetapawar0705 | github.com/ShwetaPawar0705

SUMMARY

AI/ML engineer with 9 months of internship experience in computer vision and applied ML, plus hands-on LLM work (RAG, LangChain/LangGraph agents). Cut manual video-annotation effort by 60% with a zero-shot Grounding DINO + SAM pipeline and reduced model training time by 20% through systematic hyperparameter tuning. Seeking a junior/associate role building LLM, agentic AI or computer vision applications.

EXPERIENCE

Machine Learning Intern — Ceryle Innovative Technologies Aug 2025 – Nov 2025

- Reduced average model training time by 20% and improved accuracy through systematic hyperparameter tuning (GridSearchCV, Optuna) across ML pipelines.
- Benchmarked YOLOv8 against Faster R-CNN on custom datasets and delivered comparison reports used by the team for model selection.
- Built feature engineering (normalization, encoding, dimensionality reduction) and evaluation pipelines (F1, AUC-ROC, confusion matrices) for reproducible experiments.

Data Science Intern — ACL Digital Oct 2024 – Apr 2025

- Built a zero-shot auto-labelling pipeline (Grounding DINO + SAM) that cut manual video annotation effort by 60% with no task-specific retraining.
- Designed video segmentation workflows in Python that convert frame-level model outputs into structured annotations for downstream training.
- Iterated on text prompts for text-driven detection, improving annotation precision on domain-specific video datasets, validated with cross-functional teams.

PROJECTS

RAG Q&A API | *Python, FastAPI, ChromaDB, Llama 3.1, Streamlit* | github.com/ShwetaPawar0705/RAG-Q-A

- Built a document Q&A system: FastAPI backend with ingest, query and reset endpoints, ChromaDB vector store, all-MiniLM-L6-v2 embeddings, Groq-hosted Llama 3.1 for generation, Streamlit frontend.
- Structured the codebase with Git branching and a repeatable local development workflow.

Agentic Mental Health Mapper | *LangChain, LangGraph, Streamlit* | github.com/ShwetaPawar0705/MentalHealth_State_Mapper

- Built a multi-tool LangChain/LangGraph agent that pulls data, runs sentiment analysis and generates region-level mental health risk reports.
- Shipped a Streamlit dashboard with filters for region, time window and sentiment.

Mini-Transformer: LLM from Scratch | *Python, PyTorch, Streamlit* | github.com/ShwetaPawar0705/LLM

- Implemented a decoder-only Transformer in PyTorch (multi-head self-attention, causal masking, positional embeddings, residual connections, layer norm); trained 6 layers / 384-dim / 6 heads / 256-token context on TinyShakespeare with AdamW, dropout and a 90/10 train-validation split.
- Built a Streamlit text-generation app with adjustable temperature and a CLI inference script; used CUDA acceleration where available.

TECHNICAL SKILLS

LLM & Agentic AI: LangChain, LangGraph, RAG, Prompt Engineering, OpenAI/Claude/Gemini APIs
ML/DL: PyTorch, TensorFlow, Keras, Scikit-learn, HuggingFace Transformers, XGBoost
Computer Vision: OpenCV, MediaPipe, YOLOv8, Faster R-CNN, Grounding DINO, SAM
Languages & Tools: Python, SQL, FastAPI, Streamlit, Docker, Git/GitHub, Google Cloud (Gen AI Academy labs)

CERTIFICATIONS

- Google Cloud Gen AI Academy — APAC Edition (Google Cloud x Hack2Skill): completed all 3 cohorts with a certificate for each.

EDUCATION

PG Diploma in Artificial Intelligence — CDAC 2024 | 73.07%

B.E., Electronics and Telecommunication Engineering — University of Mumbai 2023 | CGPA 8.56/10